MultiVae: A Python package for Multimodal Variational Autoencoders on Partial Datasets

Agathe Senellart[1,2,3*¶]**, Clément Chadebec**[1,2,3*]**, and Stéphanie Allassonnière**[1,2,3]

1 Université de Paris-Cité **2** Inria **3** Inserm ¶ Corresponding author * These authors contributed equally.

Summary

In recent years, there has been a major boom in the development of multimodal machine learning models. Among open topics, representation (fusion) and generation of multimodal data are very active fields of research. Recently, Multimodal Variational Autoencoders (VAEs) have been attracting growing interest for both tasks, thanks to their versatility, scalability, and interpretability as probabilistic latent variable models. They are also particularly interesting models in the *partially observed* setting, as most of them can be trained even with missing data. This last point makes them particularly interesting for the medical field, where missing data are commonplace (Antelmi et al., 2019; Lawry Aguila et al., 2023).

We present MultiVae, an open-source Python library for bringing together unified implementations of multimodal VAEs. It has been designed for easy, customizable use of these models on fully or partially observed data. This library also facilitates the development and benchmarking of new algorithms by integrating several benchmark datasets, a variety of evaluation metrics and tools for monitoring and sharing models.

Multimodal Variational Autoencoders

In Multimodal Machine Learning, two goals are generally targeted: (1) Learn a shared representation from multiple modalities; (2) Learn to generate one missing modality given the ones that are available.

Multimodal Variational Autoencoders aim at solving both issues at the same time. These models learn a latent representation z of all modalities in a lower dimensional common space and learn to *decode* z to generate any modality.

Let $X = (x_1, x_2, ... x_M)$ contain M modalities. In the VAE setting, we suppose that the generative process behind the observed data is the following:

$$z \sim p(z) \qquad \text{and} \qquad \forall 1 \leq i \leq M, x_i | z \sim p_\theta(x_i | z), \qquad (1)$$

where $p(z)$ is a prior distribution that is often fixed, and $p_\theta(x_i|z)$ are called *decoders* and are parameterized by neural network. Typically, $p_\theta(x_i|z) = \mathcal{N}(x_i; \mu_\theta(z), \sigma_\theta(z))$ where $\mu_\theta, \sigma_\theta$ are neural networks. We aim to learn these *decoders* that translate z into the high dimensional data x_i. At the same time, we aim to learn an *encoder* $q_\phi(z|X)$ that map observations to the latent space. $q_\phi(z|X)$ is also parameterized by a neural network. Derived from variational inference (Kingma & Welling, 2014), the VAE objective writes:

$$\mathcal{L}(X) = \mathbb{E}_{q_\phi(z|X)} \left(\sum_i \ln(p_\theta(x_i|z)) \right) - KL(q_\phi(z|X)|p(z))$$

.

36　The first term is a reconstruction loss and the second term can be seen as a regularization term
37　that avoids overfitting. A typical training step of a multimodal VAE consists in encoding the
38　data with the encoder, reconstructing each modality with the decoders and taking a gradient
39　step to optimize the loss $\mathcal{L}(X)$.

40　Most multimodal VAEs differ in how they construct the encoder $q_\phi(z|X)$. In the figure below,
41　we summarize several approaches: *Aggregated models* (Shi et al., 2019; Sutter et al., 2021; Wu
42　& Goodman, 2018) use a mean or a product operation to aggregate the information coming
43　from all modalities, where *Joint models* (Senellart et al., 2023; Suzuki et al., 2016; Vedantam
44　et al., 2018) use a neural network taking all modalities as input. Finally *coordinated models*
45　(Tian & Engel, 2019; Wang et al., 2017a) use different latent spaces but add a constraint
46　term in the loss to force them to be similar.

47　　　　　　　　　　　　　　　　　　⟷ : Similarity constraint

48　In our library, we implement all these approaches in an unified and modular way.

49　Aggregated models offer a natural way of learning on incomplete datasets: for an incomplete
50　sample X, we use only the available modalities to encode the data and compute the loss that
51　we minimize during training. However, except in our library MultiVae, there does not exist an
52　implementation of these models that can be used on incomplete datasets in a straightforward
53　manner.

Data Augmentation

55　Another application of these models is Data Augmentation (DA): from sampling new latent
56　codes z and decoding them, *fully synthetic multimodal* samples can be generated to augment
57　a dataset. DA has been proven useful in many data-intensive deep learning applications
58　(Chadebec et al., 2023). In a dedicated module `multivae.samplers`, we propose different
59　ways of sampling latent codes z to further explore the generative abilities of these models.

Statement of need

61　Although multimodal VAEs have interesting applications in different fields, the lack of easy-
62　to-use and verified implementations might hinder applicative research. With MultiVae, we
63　offer unified implementations, designed to be accessible even for non-specialists. In order to
64　propose reliable implementations, we reproduced, whenever possible, a key result from the
65　original paper. Some works similar to ours have grouped together model implementations:
66　the Multimodal VAE Comparison Toolkit (Sejnova et al., 2024) includes 4 models and the
67　Pixyz(Masahiro Suzuki & Matsuo, 2023) library contains 2 multimodal models. The work

⁶⁸ closest to ours and released while we were developing our library is `multi-view-ae` ([Aguila et al., 2023](#)), which contains a dozen of models. We compare in a summarizing table below, the different features of each work. Our library complements what already exists: our API is quite different compared to previous work, the models implemented are not all the same, and for those we have in common, our implementation offers additional parameterization options. Indeed, for each model, we have made sure to offer great flexibility on parameters and to include all implementation details present in the original codes. Our library also offers additional features: **compatibility with incomplete data**, which we consider essential for real-life applications, and a range of tools dedicated to research and development of new algorithms: benchmark datasets, **metrics modules** and **samplers**, for testing and analyzing models.

List of models and features

⁷⁹ In the Table below, we list available models and features, and compare to previous work. This symbol (✓*) indicates that the implementation include additional options.

Models/ Features	Ours	(Aguila et al., 2023)	(Sejnova et al., 2024)
JMVAE (Suzuki et al., 2016)	✓*	✓	
MVAE(Wu & Goodman, 2018)	✓*	✓	✓
MMVAE(Shi et al., 2019)	✓*	✓	✓
MoPoE(Sutter et al., 2021)	✓*	✓	✓
DMVAE(Lee & Pavlovic, 2021)	✓	✓*	✓
MVTCAE(Hwang et al., 2021)	✓	✓	
MMVAE+(Palumbo et al., 2023)	✓*	✓	
CMVAE(Palumbo et al., 2024)	✓		
Nexus(Vasco et al., 2022)	✓		
CVAE(Kingma & Welling, 2014)	✓		
MHVAE(Dorent et al., 2023)	✓		
TELBO(Vedantam et al., 2018)	✓		
JNF(Senellart et al., 2023)	✓		
CRMVAE(Suzuki & Matsuo, 2023)	✓		
MCVAE(Antelmi et al., 2019)		✓	
mAAE		✓	
DVCCA(Wang et al., 2017b)		✓	
mWAE		✓	
mmJSD(Sutter et al., 2020)		✓	
gPoE(Lawry Aguila et al., 2023)		✓	
Support of Incomplete datasets	✓		
GMM Sampler	✓		
MAF Sampler, IAF Sampler	✓		
Metrics: Likelihood, Coherences, FIDs, Reconstruction, Clustering	✓		
Benchmark Datasets	✓		✓
Model sharing via Hugging Face	✓		

⁸¹ An important difference in our user-interface, is that we handle all training and model parameters within python dataclasses while ([Aguila et al., 2023](#); [Sejnova et al., 2024](#)) uses independant YAML configuration files.

Code quality and documentation

Our code is available on Github (https://github.com/AgatheSenellart/MultiVae) and Pypi and we provide a full online documentation at (https://multivae.readthedocs.io/). The main features are illustrated through **tutorials** made available either as notebooks or scripts allowing users to get started easily. To further showcase how to use our library for research applications, we provide detailed case studies in the documentation.

Acknowledgements

References

Aguila, A. L., Jayme, A., Montaña-Brown, N., Heuveline, V., & Altmann, A. (2023). Multi-view-AE: A python package for multi-view autoencoder models. *Journal of Open Source Software*, *8*(85), 5093. https://doi.org/10.21105/joss.05093

Antelmi, L., Ayache, N., Robert, P., & Lorenzi, M. (2019). *Sparse multi-channel variational autoencoder for the joint analysis of heterogeneous data*. *97*, 302–311. https://proceedings.mlr.press/v97/antelmi19a.html

Chadebec, C., Thibeau-Sutre, E., Burgos, N., & Allassonnière, S. (2023). Data augmentation in high dimensional low sample size setting using a geometry-based variational autoencoder. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, *45*(3), 2879–2896. https://doi.org/10.1109/TPAMI.2022.3185773

Dorent, R., Haouchine, N., Kogl, F., Joutard, S., Juvekar, P., Torio, E., Golby, A. J., Ourselin, S., Frisken, S., Vercauteren, T., Kapur, T., & Wells, W. M. (2023). Unified brain MR-ultrasound synthesis using multi-modal hierarchical representations. In *Medical image computing and computer assisted intervention – MICCAI 2023* (pp. 448–458). Springer Nature Switzerland. https://doi.org/10.1007/978-3-031-43999-5_43

Hwang, H., Kim, G.-H., Hong, S., & Kim, K.-E. (2021). Multi-view representation learning via total correlation objective. *Advances in Neural Information Processing Systems*, *34*, 12194–12207.

Kingma, D. P., & Welling, M. (2014). *Auto-Encoding Variational Bayes*. arXiv. http://arxiv.org/abs/1312.6114

Lawry Aguila, A., Chapman, J., & Altmann, A. (2023). *Multi-modal variational autoencoders for normative modelling across multiple imaging modalities*. 425–434. https://doi.org/10.1007/978-3-031-43907-0_41

Lee, M., & Pavlovic, V. (2021). *Private-shared disentangled multimodal VAE for learning of latent representations*. 1692–1700. https://doi.org/10.1109/CVPRW53098.2021.00185

Masahiro Suzuki, T. K., & Matsuo, Y. (2023). Pixyz: A python library for developing deep generative models. In *Advanced Robotics* (No. 0; Vol. 0, pp. 1–16). Taylor & Francis. https://doi.org/10.1080/01691864.2023.2244568

Palumbo, E., Daunhawer, I., & Vogt, J. E. (2023). *MMVAE+: ENHANCING THE GENERATIVE QUALITY OF MULTIMODAL VAES WITHOUT COMPROMISES*.

Palumbo, E., Manduchi, L., Laguna, S., Chopard, D., & Vogt, J. E. (2024). *Deep generative clustering with multimodal diffusion variational autoencoders*. https://openreview.net/forum?id=k5THrhXDV3

127 Sejnova, G., Vavrecka, M., Stepanova, K., & Taniguchi, T. (2024). *Benchmarking multimodal*
128 *variational autoencoders: CdSprites+ dataset and toolkit*. https://arxiv.org/abs/2209.
129 03048

130 Senellart, A., Chadebec, C., & Allassonnière, S. (2023). Improving multimodal joint varia-
131 tional autoencoders through normalizing flows and correlation analysis. *arXiv Preprint*
132 *arXiv:2305.11832*.

133 Shi, Y., Siddharth, N., Paige, B., & Torr, P. H. S. (2019). Variational Mixture-of-Experts
134 Autoencoders for Multi-Modal Deep Generative Models. *arXiv:1911.03393 [Cs, Stat]*.
135 http://arxiv.org/abs/1911.03393

136 Sutter, T. M., Daunhawer, I., & Vogt, J. E. (2020). Multimodal generative learning utilizing
137 jensen-shannon-divergence. *CoRR*, *abs/2006.08242*. https://arxiv.org/abs/2006.08242

138 Sutter, T. M., Daunhawer, I., & Vogt, J. E. (2021). Generalized Multimodal ELBO. *ICLR*.

139 Suzuki, M., & Matsuo, Y. (2023). *Mitigating the limitations of multimodal VAEs with*
140 *coordination-based approach*. https://openreview.net/forum?id=Rn8u4MYgeNJ

141 Suzuki, M., Nakayama, K., & Matsuo, Y. (2016). Joint Multimodal Learning with Deep
142 Generative Models. *arXiv:1611.01891 [Cs, Stat]*. http://arxiv.org/abs/1611.01891

143 Tian, Y., & Engel, J. (2019). Latent Translation: Crossing Modalities by Bridging Generative
144 Models. *ArXiv*.

145 Vasco, M., Yin, H., Melo, F. S., & Paiva, A. (2022). Leveraging hierarchy in multimodal
146 generative models for effective cross-modality inference. *Neural Networks*, *146*, 238–255.

147 Vedantam, R., Fischer, I., Huang, J., & Murphy, K. (2018). Generative Models of Visually
148 Grounded Imagination. *arXiv:1705.10762 [Cs, Stat]*. http://arxiv.org/abs/1705.10762

149 Wang, W., Yan, X., Lee, H., & Livescu, K. (2017b). *Deep Variational Canonical Correlation*
150 *Analysis*. arXiv. https://doi.org/10.48550/arXiv.1610.03454

151 Wang, W., Yan, X., Lee, H., & Livescu, K. (2017a). *Deep Variational Canonical Correlation*
152 *Analysis*. arXiv. https://doi.org/10.48550/arXiv.1610.03454

153 Wu, M., & Goodman, N. (2018). Multimodal Generative Models for Scalable Weakly-Supervised
154 Learning. *Advances in Neural Information Processing Systems*, *31*. https://proceedings.
155 neurips.cc/paper/2018/hash/1102a326d5f7c9e04fc3c89d0ede88c9-Abstract.html